

The Self-Pruning Neural Network

Dynamic Architectural Adaptation via Learnable Weight Gating & Sparsity Regularization

Case Study: Tredence Analytics AI Task | Dataset: CIFAR-10

1. Executive Summary & Problem Preamble

Deploying deep neural networks on edge hardware is constrained by memory and computational budgets. Traditional post-training pruning techniques remove unneeded weights after training. This case study implements a **Self-Pruning Neural Network** that learns which of its own weights are redundant *during training*. Each weight parameter w_{ij} is associated with a learnable scalar parameter g_{ij} (gate_scores). By optimizing a total loss function combining Classification Loss (Cross-Entropy) and an L_1 Sparsity Penalty on the Sigmoid gate outputs, the network dynamically suppresses superfluous weights, adapting its architecture on the fly.

2. Custom Architectural Design: PrunableLinear Module

The custom PrunableLinear(in_features, out_features) layer replaces standard linear layers. It defines three parameter tensors: weight, bias, and gate_scores (matching the exact shape of the weight tensor).

Forward Pass Implementation:

- Gate Values: $\text{gates} = \sigma(\text{gate_scores}) \in (0, 1)$
- Pruned Weights: $\text{pruned_weights} = \text{weight} * \text{gates}$
- Layer Output: $y = x @ \text{pruned_weights}^T + \text{bias}$

3. Sparsity Regularization Loss Mechanics

Loss Formulation:

Total Loss = ClassificationLoss (Cross-Entropy) + λ * SparsityLoss

Where **SparsityLoss** is the L_1 norm of all gate values across all PrunableLinear layers: $\text{SparsityLoss} = \sum \sigma(g_{ij})$.

Why L1 Penalty Encourages Sparsity:

- Positivity of Sigmoid Outputs:** Since $\sigma(g_{ij}) \in (0, 1)$ is strictly positive, the L1 norm reduces to the sum of gate values.
- Constant Penalty Pressure:** The derivative with respect to parameter g_{ij} is $\partial(\text{SparsityLoss}) / \partial g_{ij} = \lambda * \sigma(g_{ij})(1 - \sigma(g_{ij}))$. When a connection is non-essential for reducing classification loss, this gradient continuously pushes $g_{ij} \rightarrow -\infty$, forcing $\sigma(g_{ij}) \rightarrow 0$ and zeroing out the weight.

4. Empirical Results Summary Table

Below is the summary of experimental results evaluated across multiple regularization factors (λ) from the codebase execution:


Lambda (λ)	Test Accuracy (%)	Sparsity Level (%) [Gates < 0.01]
1e-05	29.59%	0.000293%
0.0001	29.22%	0.000879%
0.0005	29.46%	0.002696%

0.001	31.27%	0.004278%
0.005	31.69%	0.014943%
0.01	30.64%	0.020100%
0.05	27.49%	0.040199%

5. Final Gate Distribution Analysis

A successful self-pruning run produces a bimodal distribution of gate values: a distinct spike at 0 (pruned weights) and a cluster of values away from 0 (active weights).

Figure 1: Distribution of Final Gate Values



6. Architectural Review & Observations

- Layer Signature & Forward Pass:** PrunableLinear perfectly conforms to Part 1 of the case study. Gradients flow seamlessly to both weight and gate_scores.
- Sparsity Loss Integration:** The L1 penalty $\sum \sigma(g_{ij})$ is computed and scaled correctly in the training loop.
- Convergence Note:** Because gate_scores start initialized around 0 (std 1), reaching $\sigma(g_{ij}) < 0.01$ requires $g_{ij} < -4.6$. Extending the training duration (epochs) or adjusting parameter-specific learning rates for gate_scores accelerates the emergence of high-sparsity bimodal spikes at 0.